

# Introduction to Parallel Computer Architecture

## Gaussian Elimination using OpenMP

Instructor: Prof. Naga Kandasamy, ECE Department, Drexel University

October 24, 2015

The assignment is due on November 1, 2015. You may work on this assignment in a team of up to two people.

Consider the problem of solving a system of linear equations of the form

$$\begin{array}{cccccc} a_{0,0}x_0 & + a_{0,1}x_1 & + \cdots & + a_{0,n-1}x_{n-1} & = b_0, \\ a_{1,0}x_0 & + a_{1,1}x_1 & + \cdots & + a_{1,n-1}x_{n-1} & = b_1, \\ \cdot & \cdot & & \cdot & \cdot \\ \cdot & \cdot & & \cdot & \cdot \\ a_{n-1,0}x_0 & + a_{n-1,1}x_1 & + \cdots & + a_{n-1,n-1}x_{n-1} & = b_{n-1}. \end{array}$$

In matrix notation, the above system is written as  $Ax = b$  where  $A$  is a dense  $n \times n$  matrix of coefficients such that  $A[i, j] = a_{i,j}$ ,  $b$  is an  $n \times 1$  vector  $[b_0, b_1, \dots, b_{n-1}]^T$ , and  $x$  is the desired solution vector  $[x_0, x_1, \dots, x_{n-1}]^T$ . From here on, we will denote the matrix elements  $a_{i,j}$  and  $x_i$  by  $A[i, j]$  and  $x[i]$ , respectively. A system of equations  $Ax = b$  is usually solved in two stages. First, through a set of algebraic manipulations, the original system of equations is reduced to an upper triangular system of the form

$$\begin{array}{cccccc} x_0 & + u_{0,1}x_1 & + u_{0,2}x_2 & + \cdots & + u_{0,n-1}x_{n-1} & = y_0, \\ & x_1 & + u_{1,2}x_2 & + \cdots & + u_{1,n-1}x_{n-1} & = y_1, \\ & & & & \cdot & \cdot \\ & & & & \cdot & \cdot \\ & & & & x_{n-1} & = y_{n-1}. \end{array}$$

We write the above system as  $Ux = y$ , where  $U$  is a unit upper-triangular matrix, that is, one where the subdiagonal entries are zero and all principal diagonal entries are equal to one. More formally,  $U[i, j] = 0$  if  $i > j$ , otherwise  $U[i, j] = u_{i,j}$ , and furthermore,  $U[i, i] = 1$  for  $0 \leq i < n$ . In the second stage of solving a system of linear equations, the upper-triangular system is solved for the variables in reverse order, from  $x[n-1]$  to  $x[0]$  using a procedure called back-substitution.

---

```

1: procedure GAUSS_ELIMINATE( $A, b, y$ )
2:   int  $i, j, k$ ;
3:   for  $k := 0$  to  $n - 1$  do
4:     for  $j := k + 1$  to  $n - 1$  do
5:        $A[k, j] := A[k, j] / A[k, k]$ ;    /* Division step. */
6:     end for
7:      $y[k] := b[k] / A[k, k]$ ;
8:      $A[k, k] := 1$ ;
9:     for  $i := k + 1$  to  $n - 1$  do
10:      for  $j := k + 1$  to  $n - 1$  do
11:         $A[i, j] := A[i, j] - A[i, k] \times A[k, j]$ ;    /* Elimination step. */
12:      end for
13:       $b[i] := b[i] - A[i, k] \times y[k]$ ;
14:       $A[i, k] := 0$ ;
15:    end for
16:  end for

```

---

A serial implementation of a simple Gaussian elimination algorithm is shown below. The algorithm converts the system of linear equations  $Ax = b$  into a unit upper-triangular system  $Ux = y$ . We assume that the matrix  $u$  shares storage with  $A$  and overwrites the upper-triangular portion of  $A$ . So, the element  $A[k, j]$  computed in line 5 of the code is actually  $U[k, j]$ . Similarly, the element  $A[k, k]$  that is equated to 1 in line 8 is  $U[k, k]$ . Also, our program assumes that  $A[k, k] \neq 0$  when it is used as a divisor in lines 5 and 7. So, our implementation is numerically unstable, though it should not be a concern for this assignment. For  $k$  ranging from 0 to  $n - 1$ , the Gaussian elimination procedure systematically eliminates the variable  $x[k]$  from equations  $k + 1$  to  $n - 1$  so that the matrix of coefficients becomes upper-triangular. In the  $k^{\text{th}}$  iteration of the outer loop (line 3), an appropriate multiple of the  $k^{\text{th}}$  equation is subtracted from each of the equations  $k + 1$  to  $n - 1$ .

This problem asks you optimize the performance of GAUSS\_ELIMINATE using SSE. The program given to you accepts no arguments. The CPU computes the reference solution which is compared with the result provided by the SSE-based implementation. If the solutions match, the application will print out “Test PASSED” to the screen before exiting.

Edit the `gauss_eliminate_using_sse()` function in `gauss_eliminate.c` to complete the functionality of Gaussian elimination using SSE. The size of the matrix  $A$  is guaranteed to be  $2048 \times 2048$ . Build the code as follows:

Submit all of the files needed to run your code as a single zip file via BBLearn. Also include a brief report describing how you designed your SSE-optimized code, using code or pseudocode to help the discussion, and the amount of speedup obtained over the reference version.